

# Session 7:

# EXPLORING APACHE PIG

## Assignment

### Task 1

Write a program to implement wordcount using Pig.

```
A = load '/test.txt';
B = foreach A generate flatten(TOKENIZE((chararray)$0)) as word;
C = group B by word;
D = foreach C generate group, COUNT(B);
dump D;
```

```
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -cat /test.txt
18/08/19 14:22:07 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Hello World.Have a nice day.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

```
2018-08-19 14:23:12,393 [main] WARN org.apache.pig.PigServer - ATS is disabled since yarn.timeline-service.enabled set to false
grunt> A = load '/test.txt';
2018-08-19 14:24:05,719 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> B = foreach A generate flatten(TOKENIZE((chararray)$0)) as word;
2018-08-19 14:24:24,608 [main] INFO org.apache.pig.impl.util.SpillableMemoryManager - Selected heap (Tenured Gen) of size 699072512 to monitor. collectionUsageThreshold = 489350752, usageThreshold = 489350752
grunt> C = group B by word;
grunt> D = foreach C generate group, COUNT(B);
grunt> dump D;
2018-08-19 14:25:10,022 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: GROUP_BY
2018-08-19 14:25:10,126 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-08-19 14:25:52,317 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code
2018-08-19 14:25:52,367 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 14:25:52,367 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(a,1)
(day,,1)
(nice,,1)
>Hello,1)
(World.Have,1)
grunt>
```

### Task 2

We have employee\_details and employee\_expenses files. Use local mode while running Pig and write Pig Latin script to get below results:

(a) Top 5 employees (employee id and employee name) with highest rating. (In case two

employees have same rating, employee with name coming first in dictionary should get preference)

```
(world.Have,1)
grunt> emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int, emp_name:chararray, emp_salary:int,emp_rating:int);
2018-08-19 14:52:39,006 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> dump emp;
2018-08-19 14:53:05,012 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2018-08-19 14:53:05,072 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
```

```
2018-08-19 14:53:37,028 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh,20000,1)
(102,Shahrukh,10000,2)
(103,Akshay,11000,3)
(104,Anubhav,5000,4)
(105,Pawan,2500,5)
(106,Aamir,25000,1)
(107,Salman,17500,2)
(108,Ranbir,14000,3)
(109,Katrina,1000,4)
(110,Priyanka,2000,5)
(111,Tushar,500,1)
(112,Ajay,5000,2)
(113,Jubeen,1000,1)
(114,Madhuri,2000,2)
grunt>
```

```
Details at logfile: /home/acadgild/pig_1534868789659.log
grunt> emp_expenses= LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,expenses:int);
2018-08-19 14:57:08,090 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> dump emp_expenses;
2018-08-19 14:57:16,903 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2018-08-19 14:57:16,973 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
```

```
2018-08-19 14:57:43,663 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,200)
(102,100)
(110,400)
(114,200)
(119,200)
(105,100)
(101,100)
(104,300)
(102,400)
grunt>
```

Pig Query:

```
emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int, emp_name:chararray, emp_salary:int,emp_rating:int);
```

```
rating = order emp by emp_rating DESC;
```

```
Result = LIMIT rating 5;
```

Dump Result;

```
(102,400)
grunt> emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int, emp_name:chararray, emp_salary:int,emp_rating:int);
2018-08-19 14:59:25,966 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> rating = order emp by emp_rating DESC;
grunt> Result = LIMIT rating 5;
grunt> Dump Result;
2018-08-19 15:00:02,106 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: ORDER BY,LIMIT
2018-08-19 15:00:02,253 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-08-19 15:00:02,255 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code
2018-08-19 15:00:02,256 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMajorKeyPrune, ConstantCalculator, GroupByConstParallelSetter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, PartitionFilterOptimizer, PredicatePushdownOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter]}
2018-08-19 15:00:02,275 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold
```

```

2018-08-19 15:02:23,703 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 15:02:23,703 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka,2000,5)
(105,Pawan,2500,5)
(109,Katrina,1000,4)
(104,Anubhav,5000,4)
(108,Ranbir,14000,3)
grunt> █

```

b) Top 3 employees (employee id and employee name) with highest salary, whose employee id is an odd number. (In case two employees have same salary, employee with name coming first in dictionary should get preference)

Pig Query:

```

emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
emp_name:chararray, emp_salary:int,emp_rating:int);
emp_sal_name = order emp by emp_salary desc;
emp_sal_id = FILTER emp_sal_name by emp_id%2==1;
emp_final = FOREACH emp_sal_id generate emp_id,emp_name;
emp_final_limit = LIMIT emp_final 3;

```

```

Details at logfile: /home/acadgild/pig_1534668789659.log
grunt> emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
>> emp_name:chararray, emp_salary:int,emp_rating:int);
2018-08-19 15:08:34,571 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
grunt> emp_sal_name = order emp by emp_salary desc;
grunt> emp_sal_id = FILTER emp_sal_name by emp_id%2==1;
grunt> emp_final = FOREACH emp_sal_id generate emp_id,emp_name;
grunt> emp_final_limit = LIMIT emp_final 3;
grunt> dump emp_final_limit;
2018-08-19 15:09:43,283 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: ORDER_BY,FILTER,LIMIT
2018-08-19 15:09:43,323 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS

```

```

2018-08-19 15:11:58,908 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 15:11:58,908 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh)
(107,Salman)
(103,Akshay)
grunt> █

```

(c) Employee (employee id and employee name) with maximum expense (In case two employees have same expense, employee with name coming first in dictionary should get preference)

Pig Query:

```

emp = LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
emp_name:chararray, emp_salary:int);
emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
expenses:int);
Joinempempexpense = join emp by emp_id, emp_expenses by emp_id;
maxexpense = ORDER Joinempempexpense by emp_expenses::expenses desc;
Limitmaxepnse = LIMIT maxexpense 1;
Limitmaxexpensefinal = foreach Limitmaxepnse generate emp::emp_id,emp::emp_name;

```

dump Limitmaxexpensefinal;

```
grunt> emp = LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
>> emp_name:chararray, emp_salary:int);
2018-08-19 15:15:04,595 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
grunt> emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
>> expenses:int);
2018-08-19 15:15:23,428 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
grunt> Joinempempexpense = join emp by emp_id, emp_expenses by emp_id;
grunt> maxexpense = ORDER Joinempempexpense by emp_expenses::expenses desc;
grunt> Limitmaxexpnse = LIMIT maxexpense 1;
grunt> Limitmaxexpensefinal = foreach Limitmaxexpnse generate emp::emp_id,emp::emp_name;
grunt> dump Limitmaxexpensefinal;
2018-08-19 15:16:41,911 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: HASH_JOIN,ORDER_BY,LIMIT
2018-08-19 15:19:16,076 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 15:19:16,076 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(110,Priyanka)
grunt>
```

(d) List of employees (employee id and employee name) having entries in employee\_expenses file.

Pig Query:

```
emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
emp_name:chararray, emp_salary:int,emp_rating:int);
emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
expenses:int);
emp_with_exp = JOIN emp BY emp_id, emp_expenses BY emp_id;
emp_with_exp_data = FOREACH emp_with_exp GENERATE emp::emp_id,
emp::emp_name;
emp_with_exp_distinct_data = DISTINCT emp_with_exp_data;
dump emp_with_exp_distinct_data;
```

```
(110,Priyanka)
grunt> emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
>> emp_name:chararray, emp_salary:int,emp_rating:int);
2018-08-19 15:21:43,364 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
grunt> emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
>> expenses:int);
2018-08-19 15:21:43,637 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
grunt> emp_with_exp = JOIN emp BY emp_id, emp_expenses BY emp_id;
grunt> emp_with_exp_data = FOREACH emp_with_exp GENERATE emp::emp_id,
>> emp::emp_name;
grunt> emp_with_exp_distinct_data = DISTINCT emp_with_exp_data;
grunt> dump emp_with_exp_distinct_data;
2018-08-19 15:21:44,929 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: HASH_JOIN,DISTINCT
2018-08-19 15:21:45,004 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.de
faultFS
```

```
2018-08-19 15:23:02,076 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 15:23:02,076 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(101,Amitabh)
(102,Shahrukh)
(104,Anubhav)
(105,Pawan)
(110,Priyanka)
(114,Madhuri)
grunt>
```

(e) List of employees (employee id and employee name) having no entry in employee\_expenses

file.

Pig Query:

```
emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
emp_name:chararray, emp_salary:int,emp_rating:int);
emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
expenses:int);
emp_without_exp = JOIN emp BY emp_id LEFT OUTER, emp_expenses BY emp_id;
emp_without_exp_filter = FILTER emp_without_exp BY emp_expenses::emp_id is null;
emp_without_exp_filter_data = FOREACH emp_without_exp_filter GENERATE
emp::emp_id, emp::emp_name;
DUMP emp_without_exp_filter_data;
```

```
(114,Madhuri)
grunt> emp= LOAD '/hadoopdata/pig/employee_details.txt' USING PigStorage(',') AS (emp_id:int,
>> emp_name:chararray, emp_salary:int,emp_rating:int);
2018-08-19 15:25:34,714 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> emp_expenses = LOAD '/hadoopdata/pig/employee_expenses.txt' AS (emp_id:int,
>> expenses:int);
2018-08-19 15:25:34,964 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
grunt> emp_without_exp = JOIN emp BY emp_id LEFT OUTER, emp_expenses BY emp_id;
grunt> emp_without_exp_filter = FILTER emp_without_exp BY emp_expenses::emp_id is null;
grunt> emp_without_exp_filter_data = FOREACH emp_without_exp_filter GENERATE
grunt> emp::emp_id, emp::emp_name;
grunt> DUMP emp_without_exp_filter_data;
2018-08-19 15:25:36,076 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: HASH_JOIN,FILTER
2018-08-19 15:25:36,123 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2018-08-19 15:25:36,123 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code
2018-08-19 15:25:36,124 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - (RULES_ENABLED) [AddForEach, ColumnM...
```

```
2018-08-19 15:26:17,551 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2018-08-19 15:26:17,551 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(103,Akshay)
(106,Aamir)
(107,Salman)
(108,Ranbir)
(109,Katrina)
(111,Tushar)
(112,Ajay)
(113,Jubeen)
grunt> █
```

### Task 3

Implement the use case present in below blog link and share the complete steps along with screenshot(s) from your end.

<https://acadgild.com/blog/aviation-data-analysis-using-apache-pig/>

## Problem Statement 1

Find out the top 5 most visited destinations.

REGISTER '/hadoopdata/pig/piggybank.jar';

A = load '/hadoopdata/pig/DelayedFlights.csv' USING

org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO\_MULTILINE', 'UNIX', 'SKIP\_INPUT\_HEADER');

B = foreach A generate (int)\$1 as year, (int)\$10 as flight\_num, (chararray)\$17 as origin, (chararray) \$18 as dest;

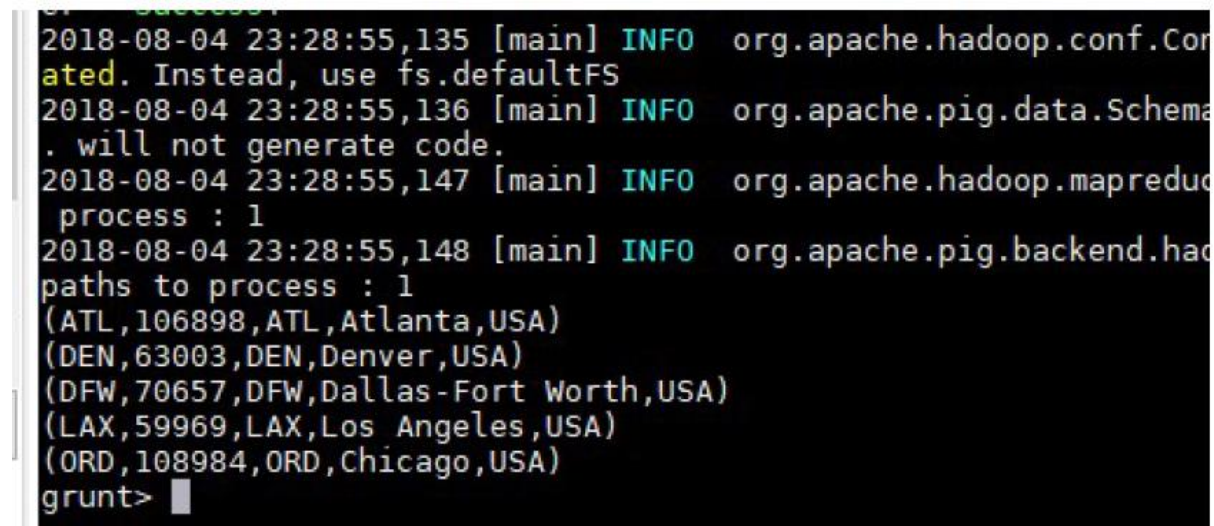
C = filter B by dest is not null;

```

D = group C by dest;
E = foreach D generate group, COUNT(C.dest);
F = order E by $1 DESC;
Result = LIMIT F 5;
A1 = load 'hadoopdata/pig/airports.csv' USING
org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_I
NPUT_HEADER');

A2 = foreach A1 generate (chararray)$0 as dest, (chararray)$2 as city, (chararray)$4 as
country;
joined_table = join Result by $0, A2 by dest;
dump joined_table;

```



```

2018-08-04 23:28:55,135 [main] INFO org.apache.hadoop.conf.Configuration: Using
ated. Instead, use fs.defaultFS
2018-08-04 23:28:55,136 [main] INFO org.apache.pig.data.SchemaTool: SchemaTool
. will not generate code.
2018-08-04 23:28:55,147 [main] INFO org.apache.hadoop.mapreduce.lib.input.TextInputFormat:
process : 1
2018-08-04 23:28:55,148 [main] INFO org.apache.pig.backend.hadoop.mapreduce.Execu
paths to process : 1
(ATL,106898,ATL,Atlanta,USA)
(DEN,63003,DEN,Denver,USA)
(DFW,70657,DFW,Dallas-Fort Worth,USA)
(LAX,59969,LAX,Los Angeles,USA)
(ORD,108984,ORD,Chicago,USA)
grunt>

```

## Problem Statement 2

Which month has seen the most number of cancellations due to bad weather?

Source code

```

REGISTER '/hadoopdata/pig/piggybank.jar';
A = load '/hadoopdata/pig/DelayedFlights.csv' USING
org.apache.pig.piggybank.storage.CSVExcelStorage('','NO_MULTILINE','UNIX','SKIP_I
NPUT_HEADER');
B = foreach A generate (int)$2 as month,(int)$10 as flight_num,(int)$22 as
cancelled,(chararray)$23 as cancel_code;
C = filter B by cancelled == 1 AND cancel_code == 'B';
D = group C by month;
E = foreach D generate group, COUNT(C.cancelled);
F = order E by $1 DESC;
Result = limit F 1;

```



dump Result;

```
ted. FinalApplicationStatus=SUCCEEDED. Redirecting to job history server
2018-08-04 23:35:28,852 [main] INFO org.apache.pig.backend.hadoop.execution
er - Success!
2018-08-04 23:35:28,852 [main] INFO org.apache.hadoop.conf.Configuration
ated. Instead, use fs.defaultFS
2018-08-04 23:35:28,856 [main] INFO org.apache.pig.data.SchemaTupleBacker
. will not generate code.
2018-08-04 23:35:28,862 [main] INFO org.apache.hadoop.mapreduce.lib.input
process : 1
2018-08-04 23:35:28,862 [main] INFO org.apache.pig.backend.hadoop.execution
paths to process : 1
(12,250)
grunt> █
```

## Problem Statement 3

Top ten origins with the highest AVG departure delay

Source code

```
A = load '/hadoopdata/pig/DelayedFlights.csv' USING
org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'UNIX', 'SKIP_I
NPUT_HEADER');
B1 = foreach A generate (int)$16 as dep_delay, (chararray)$17 as origin;
C1 = filter B1 by (dep_delay is not null) AND (origin is not null);
D1 = group C1 by origin;
E1 = foreach D1 generate group, AVG(C1.dep_delay);
Result = order E1 by $1 DESC;
Top_ten = limit Result 10;
Lookup = load '/hadoopdata/pig/airports.csv' USING
org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'UNIX', 'SKIP_I
NPUT_HEADER');
Lookup1 = foreach Lookup generate (chararray)$0 as origin, (chararray)$2 as city,
(chararray)$4 as country;
Joined = join Lookup1 by origin, Top_ten by $0;
Final = foreach Joined generate $0,$1,$2,$4;
Final_Result = ORDER Final by $3 DESC;
dump Final_Result;
```

```

2018-08-04 23:41:44,624 [main] INFO org.apache.hadoop.conf.Configuration
ated. Instead, use fs.defaultFS
2018-08-04 23:41:44,624 [main] INFO org.apache.pig.data.SchemaTupleBacke
. will not generate code.
2018-08-04 23:41:44,632 [main] INFO org.apache.hadoop.mapreduce.lib.inpu
process : 1
2018-08-04 23:41:44,632 [main] INFO org.apache.pig.backend.hadoop.execut
paths to process : 1
(CMX,Hancock,USA,116.1470588235294)
(PLN,Pellston,USA,93.76190476190476)
(SPI,Springfield,USA,83.84873949579831)
(ALO,Waterloo,USA,82.2258064516129)
(MQT,NA,USA,79.55665024630542)
(ACY,Atlantic City,USA,79.3103448275862)
(MOT,Minot,USA,78.66165413533835)
(HHH,NA,USA,76.53005464480874)
(EGE,Eagle,USA,74.12891986062718)
(BGM,Binghamton,USA,73.15533980582525)
grunt> █

```

## Problem Statement 4

Which route (origin & destination) has seen the maximum diversion?

Source code

```

A = load '/hadoopdata/pig/DelayedFlights.csv' USING
org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'NO_MULTILINE', 'UNIX', 'SKIP_I
NPUT_HEADER');
B = FOREACH A GENERATE (chararray)$17 as origin, (chararray)$18 as dest, (int)$24
as diversion;
C = FILTER B BY (origin is not null) AND (dest is not null) AND (diversion == 1);

D = GROUP C by (origin, dest);
E = FOREACH D generate group, COUNT(C.diversion);
F = ORDER E BY $1 DESC;
Result = limit F 10;
dump Result;

```



```
er - Success!
2018-08-04 23:45:25,381 [main] INFO org.apache.hadoop.conf.C
ated. Instead, use fs.defaultFS
2018-08-04 23:45:25,384 [main] INFO org.apache.pig.data.Sche
. will not generate code.
2018-08-04 23:45:25,387 [main] INFO org.apache.hadoop.mapred
process : 1
2018-08-04 23:45:25,387 [main] INFO org.apache.pig.backend.h
paths to process : 1
((ORD,LGA),39)
((DAL,HOU),35)
((DFW,LGA),33)
((ATL,LGA),32)
((ORD,SNA),31)
((SLC,SUN),31)
((MIA,LGA),31)
((BUR,JFK),29)
((HRL,HOU),28)
((BUR,DFW),25)
grunt> █
```